

Task 1

```
In [1]:  #Taking housing prices dataset from kaggle to work on this task
```

```
In [2]:  import pandas as pd
import numpy as np
```


```
In [3]:  df=pd.read_csv("./Housing_train.csv")
```

```
In [4]:  df.head()
```

Out[4]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl

5 rows × 81 columns




In [5]: ▶ `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig            1460 non-null   object
11  LandSlope            1460 non-null   object
12  Neighborhood         1460 non-null   object
13  Condition1           1460 non-null   object
14  Condition2           1460 non-null   object
15  BldgType             1460 non-null   object
16  HouseStyle           1460 non-null   object
17  OverallQual          1460 non-null   int64
18  OverallCond          1460 non-null   int64
19  YearBuilt            1460 non-null   int64
20  YearRemodAdd         1460 non-null   int64
21  RoofStyle            1460 non-null   object
22  RoofMatl            1460 non-null   object
23  Exterior1st          1460 non-null   object
24  Exterior2nd          1460 non-null   object
25  MasVnrType           1452 non-null   object
26  MasVnrArea           1452 non-null   float64
27  ExterQual            1460 non-null   object
28  ExterCond            1460 non-null   object
29  Foundation           1460 non-null   object
30  BsmtQual             1423 non-null   object
31  BsmtCond             1423 non-null   object
32  BsmtExposure         1422 non-null   object
33  BsmtFinType1         1423 non-null   object
34  BsmtFinSF1           1460 non-null   int64
35  BsmtFinType2         1422 non-null   object
36  BsmtFinSF2           1460 non-null   int64
37  BsmtUnfSF            1460 non-null   int64
38  TotalBsmtSF          1460 non-null   int64
39  Heating              1460 non-null   object
40  HeatingQC            1460 non-null   object
41  CentralAir           1460 non-null   object
42  Electrical           1459 non-null   object
43  1stFlrSF             1460 non-null   int64
44  2ndFlrSF             1460 non-null   int64
45  LowQualFinSF         1460 non-null   int64
46  GrLivArea            1460 non-null   int64
47  BsmtFullBath         1460 non-null   int64
48  BsmtHalfBath         1460 non-null   int64
49  FullBath             1460 non-null   int64
50  HalfBath             1460 non-null   int64
51  BedroomAbvGr         1460 non-null   int64
```

```

52 KitchenAbvGr 1460 non-null int64
53 KitchenQual 1460 non-null object
54 TotRmsAbvGrd 1460 non-null int64
55 Functional 1460 non-null object
56 Fireplaces 1460 non-null int64
57 FireplaceQu 770 non-null object
58 GarageType 1379 non-null object
59 GarageYrBlt 1379 non-null float64
60 GarageFinish 1379 non-null object
61 GarageCars 1460 non-null int64
62 GarageArea 1460 non-null int64
63 GarageQual 1379 non-null object
64 GarageCond 1379 non-null object
65 PavedDrive 1460 non-null object
66 WoodDeckSF 1460 non-null int64
67 OpenPorchSF 1460 non-null int64
68 EnclosedPorch 1460 non-null int64
69 3SsnPorch 1460 non-null int64
70 ScreenPorch 1460 non-null int64
71 PoolArea 1460 non-null int64
72 PoolQC 7 non-null object
73 Fence 281 non-null object
74 MiscFeature 54 non-null object
75 MiscVal 1460 non-null int64
76 MoSold 1460 non-null int64
77 YrSold 1460 non-null int64
78 SaleType 1460 non-null object
79 SaleCondition 1460 non-null object
80 SalePrice 1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB


```

In [6]:  *#checking for null values*
`print(df.isnull().sum())`

```

Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea          0
...
MoSold           0
YrSold           0
SaleType         0
SaleCondition    0
SalePrice        0
Length: 81, dtype: int64

```

In [7]:  *#filling the null values with most common values*
`df.fillna(df.mode().iloc[0], inplace=True)`

```
In [8]: #standardising the values  
from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()
```

```
In [9]: #numerical values  
df.select_dtypes(include=np.number).columns.tolist()
```

```
Out[9]: ['Id',  
        'MSSubClass',  
        'LotFrontage',  
        'LotArea',  
        'OverallQual',  
        'OverallCond',  
        'YearBuilt',  
        'YearRemodAdd',  
        'MasVnrArea',  
        'BsmtFinSF1',  
        'BsmtFinSF2',  
        'BsmtUnfSF',  
        'TotalBsmtSF',  
        '1stFlrSF',  
        '2ndFlrSF',  
        'LowQualFinSF',  
        'GrLivArea',  
        'BsmtFullBath',  
        'BsmtHalfBath',  
        'FullBath',  
        'HalfBath',  
        'BedroomAbvGr',  
        'KitchenAbvGr',  
        'TotRmsAbvGrd',  
        'Fireplaces',  
        'GarageYrBlt',  
        'GarageCars',  
        'GarageArea',  
        'WoodDeckSF',  
        'OpenPorchSF',  
        'EnclosedPorch',  
        '3SsnPorch',  
        'ScreenPorch',  
        'PoolArea',  
        'MiscVal',  
        'MoSold',  
        'YrSold',  
        'SalePrice']
```

```
In [10]: num=df.select_dtypes(include=np.number).columns.tolist()  
df[num]=scaler.fit_transform(df[num])
```

In [11]: `df[num]`

Out[11]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt
0	-1.730865	0.073375	-0.146189	-0.207142	0.651479	-0.517200	1.050994
1	-1.728492	-0.872563	0.524992	-0.091886	-0.071836	2.179628	0.156734
2	-1.726120	0.073375	-0.011953	0.073480	0.651479	-0.517200	0.984752
3	-1.723747	0.309859	-0.369915	-0.096897	0.651479	-0.517200	-1.863632
4	-1.721374	0.073375	0.703973	0.375148	1.374795	-0.517200	0.951632
...
1455	1.721374	0.073375	-0.280425	-0.260560	-0.071836	-0.517200	0.918511
1456	1.723747	-0.872563	0.748718	0.266407	-0.071836	0.381743	0.222975
1457	1.726120	0.309859	-0.101443	-0.147810	0.651479	3.078570	-1.002492
1458	1.728492	-0.872563	-0.011953	-0.080160	-0.795151	0.381743	-0.704406
1459	1.730865	-0.872563	0.301265	-0.058112	-0.795151	0.381743	-0.207594

1460 rows × 38 columns

```

In [12]: #handling categorical values
from sklearn.preprocessing import LabelEncoder
categorical=df.select_dtypes(include=[object]).columns.tolist()
le=LabelEncoder()
df[categorical] = df[categorical].apply(lambda col: le.fit_transform(col))
df.head()

```

Out[12]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	Land
0	-1.730865	0.073375	3	-0.146189	-0.207142	1	0	3	
1	-1.728492	-0.872563	3	0.524992	-0.091886	1	0	3	
2	-1.726120	0.073375	3	-0.011953	0.073480	1	0	0	
3	-1.723747	0.309859	3	-0.369915	-0.096897	1	0	0	
4	-1.721374	0.073375	3	0.703973	0.375148	1	0	0	

5 rows × 81 columns

```
In [13]: #using one hot encoding
onehotencoding=pd.get_dummies(df, drop_first=True)
onehotencoding.head()
```

```
Out[13]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	Land
0	-1.730865	0.073375	3	-0.146189	-0.207142	1	0	3	
1	-1.728492	-0.872563	3	0.524992	-0.091886	1	0	3	
2	-1.726120	0.073375	3	-0.011953	0.073480	1	0	0	
3	-1.723747	0.309859	3	-0.369915	-0.096897	1	0	0	
4	-1.721374	0.073375	3	0.703973	0.375148	1	0	0	

5 rows × 81 columns

```
In [14]: #multicollinearity
#the more the value the more collinear the features are
df.corr()
```

```
Out[14]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Al
Id	1.000000	0.011156	-0.006096	-0.012497	-0.033226	0.008916	-0.0016
MSSubClass	0.011156	1.000000	0.035900	-0.349116	-0.139781	-0.024969	0.1841
MSZoning	-0.006096	0.035900	1.000000	-0.101150	-0.034452	0.087654	-0.3291
LotFrontage	-0.012497	-0.349116	-0.101150	1.000000	0.281283	-0.037078	-0.1591
LotArea	-0.033226	-0.139781	-0.034452	0.281283	1.000000	-0.197131	-0.0771
...
MoSold	0.021172	-0.013585	-0.031496	0.012785	0.001205	0.003690	-0.0211
YrSold	0.000712	-0.021407	-0.020628	0.003021	-0.014261	-0.025043	-0.0011
SaleType	0.019773	0.012464	0.097437	-0.035773	0.012292	0.014339	-0.0051
SaleCondition	-0.005806	-0.024940	0.009494	0.061393	0.034169	0.006064	-0.0281
SalePrice	-0.021917	-0.084284	-0.166872	0.329220	0.263843	0.041036	-0.0271

81 rows × 81 columns

Task 2

```
In [15]: import pandas as pd
```

```
In [16]: # Not using only 643 rows becuase the whole data set was crashing
# the system while performing label encoding
```

```
In [17]: with open('./farm-data/farm-ads.txt') as f:
        lines = f.readlines()
        len(lines)
        temp=[]
        for i in range(len(lines)-3500):
            temp.append(lines[i].split())
        farm_ad=pd.DataFrame(temp)
        farm_ad.head()
```

```
Out[17]:
```

	0	1	2	3	4	5	6	7	
0	1	ad-jerry	ad-bruckheimer	ad-chase	ad-premier	ad-sept	ad-th	ad-clip	bruckheir
1	-1	ad-rheumatoid	ad-arthritis	ad-expert	ad-tip	ad-info	ad-article	ad-treatment	ad-opt
2	-1	ad-rheumatologist	ad-anju	ad-varghese	ad-yonker	ad-ny	ad-pomona	ad-ny	ad-w
3	-1	ad-siemen	ad-water	ad-remediation	ad-water	ad-scarce	ad-resource	ad-siemen	ad-h
4	-1	ad-symptom	ad-muscle	ad-weakness	ad-genetic	ad-disease	ad-symptom	ad-include	ad-sea

5 rows × 7373 columns

```
In [18]: with open('./farm-data/farm-ads-vect.txt') as f:
        lines = f.readlines()
        len(lines)
        temp=[]
        for i in range(len(lines)-3500):
            temp.append(lines[i].split())
        farm_ad_vec=pd.DataFrame(temp)
        farm_ad_vec.head()
```

```
Out[18]:
```

	0	1	2	3	4	5	6	7	8	9	...	1737	1738	1739	17
0	1	1:1	2:1	3:1	4:1	5:1	6:1	7:1	8:1	9:1	...	None	None	None	Nc
1	-1	10:1	11:1	12:1	13:1	14:1	15:1	16:1	17:1	18:1	...	None	None	None	Nc
2	-1	29:1	31:1	35:1	101:1	131:1	252:1	272:1	280:1	291:1	...	None	None	None	Nc
3	-1	34:1	35:1	36:1	44:1	54:1	84:1	94:1	104:1	126:1	...	None	None	None	Nc
4	-1	8:1	9:1	429:1	430:1	431:1	432:1	433:1	434:1	435:1	...	None	None	None	Nc

5 rows × 1747 columns

In [19]: `data=pd.merge(farm_ad,farm_ad_vec,left_index=True,right_index=True)`
`data`

Out[19]:

	0_x	1_x	2_x	3_x	4_x	5_x	6_x	7_x
0	1	ad-jerry	ad-bruckheimer	ad-chase	ad-premier	ad-sept	ad-th	ad-clip
1	-1	ad-rheumatoid	ad-arthritis	ad-expert	ad-tip	ad-info	ad-article	ad-treatment
2	-1	ad-rheumatologist	ad-anju	ad-varghese	ad-yonker	ad-ny	ad-pomona	ad-ny
3	-1	ad-siemen	ad-water	ad-remediation	ad-water	ad-scarce	ad-resource	ad-siemer
4	-1	ad-symptom	ad-muscle	ad-weakness	ad-genetic	ad-disease	ad-symptom	ad-include
...
638	-1	ad-fibromyalgia	ad-free	ad-hopkin	ad-fibromyalgia	ad-report	ad-learn	ad-stop
639	-1	ad-infectious	ad-disease	ad-online	ad-cost	ad-conference	ad-cme	ad-credit
640	-1	ad-local	ad-sperm	ad-bank	ad-sperm	ad-donor	ad-pregnant	ad-request
641	1	ad-aoudad	ad-whitetail	ad-hunt	ad-ram	ad-white	ad-tail	ad-low
642	-1	ad-egg	ad-donation	ad-service	ad-low	ad-cost	ad-birth	ad-success

643 rows × 9120 columns

In [20]: `from sklearn.preprocessing import LabelEncoder`
`categorical=data.select_dtypes(include=[object]).columns.tolist()`
`le=LabelEncoder()`
`data[categorical] = data[categorical].apply(lambda col: le.fit_transform(c`
`# data['0_y'] = data['0_y'].apply(lambda col: le.fit_transform(col))`
`# data['0_x']=data['0_x'].astype('category').cat.codes`
`# data['0_y']=data['0_y'].astype('category').cat.codes`

```
In [21]: # data['0_x']=data['0_x'].astype('category').cat.codes
# data['0_y']=data['0_y'].astype('category').cat.codes
data['target'] = data['0_x'] + data['0_y']
data = data.drop(['0_x','0_y'],axis = 1)
data
```

C:\Users\Aaryan Agarwal\AppData\Local\Temp\ipykernel_31852\3533153650.py:
 3: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use `newframe = frame.copy()`
 data['target'] = data['0_x'] + data['0_y']

```
Out[21]:
```

	1_x	2_x	3_x	4_x	5_x	6_x	7_x	8_x	9_x	10_x	...	1738_y	1739_y	1740_y	1741
0	162	29	40	218	268	317	53	23	33	212	...	1	1	1	
1	251	10	99	295	156	12	315	182	217	363	...	1	1	1	
2	252	6	303	327	200	226	209	293	12	40	...	1	1	1	
3	263	310	232	317	260	265	277	118	159	350	...	1	1	1	
4	282	184	310	129	89	306	156	228	115	212	...	1	1	1	
...	
638	112	112	135	120	249	170	290	186	307	302	...	1	1	1	
639	152	80	195	64	61	52	61	135	98	175	...	1	1	1	
640	177	265	14	272	96	232	256	107	106	102	...	1	1	1	
641	17	313	143	237	334	307	189	105	173	152	...	1	1	1	
642	93	82	250	169	65	21	293	212	51	25	...	1	1	1	

643 rows × 9119 columns



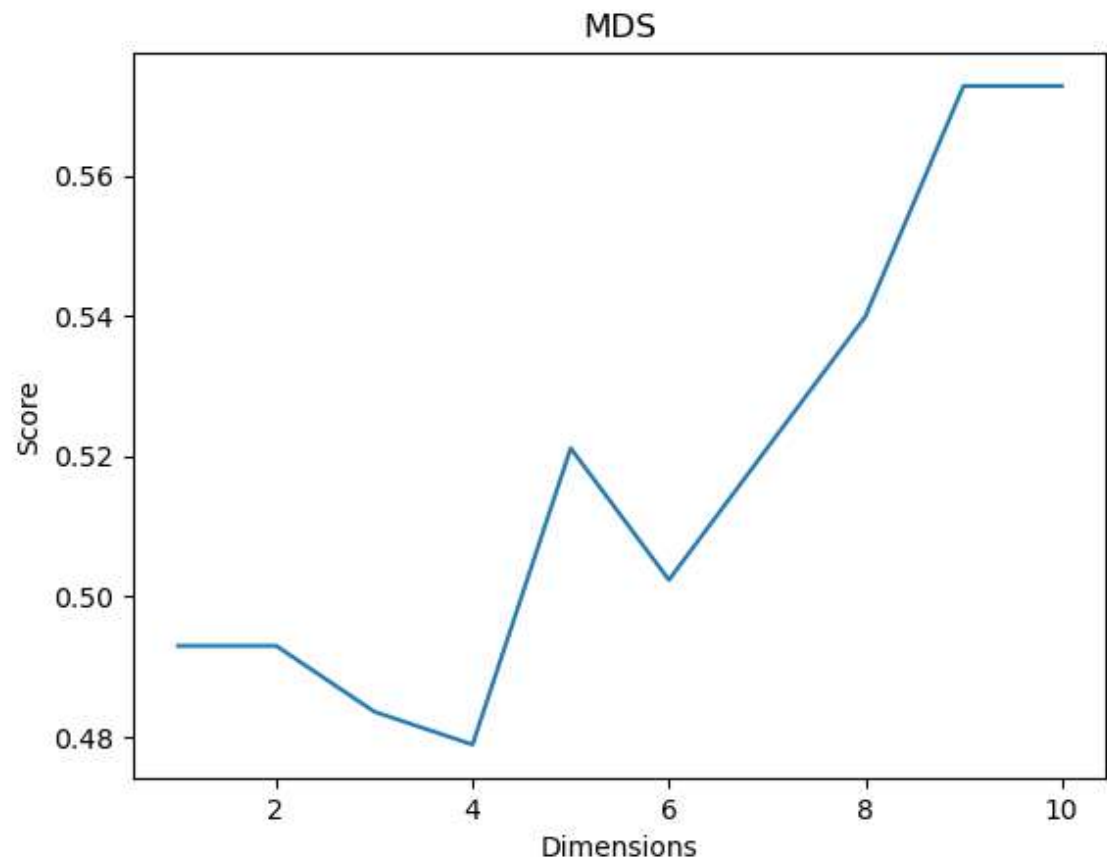
```
In [22]: ▶ #MDS
from scipy.spatial.distance import pdist,squareform
from sklearn.manifold import MDS
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

X = data.drop('target',axis = 1)
y = data['target']
scores=[]
dismat=squareform(pdist(data))

for i in range(1,11):
    mds=MDS(n_components=i,dissimilarity='precomputed')
    dim_red=mds.fit_transform(dismat)
    X_train, X_test, y_train, y_test = train_test_split(dim_red, y, test_s
    lr=LogisticRegression()
    lr.fit(X_train,y_train)
    scores.append(lr.score(X_test,y_test))
print(scores)
```

```
[0.49295774647887325, 0.49295774647887325, 0.4835680751173709, 0.47887323
94366197, 0.5211267605633803, 0.5023474178403756, 0.5211267605633803, 0.5
39906103286385, 0.5727699530516432, 0.5727699530516432]
```

```
In [23]: ▶ import matplotlib.pyplot as plt
plt.title("MDS")
plt.xlabel("Dimensions")
plt.ylabel("Score")
plt.plot(list(range(1,11)), scores)
plt.show()
```



```
In [24]: from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
X = data.drop('target',axis = 1)
y = data['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
scores=[])
for i in range(1,11):
    alg=PCA(n_components=i)
    X_train_reduced = alg.fit_transform(X_train)
    lr = LogisticRegression()
    lr.fit(X_train_reduced, y_train)
    X_test_reduced = alg.transform(X_test)
    scores.append(lr.score(X_test_reduced, y_test))
scores
```

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

warnings.warn(

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

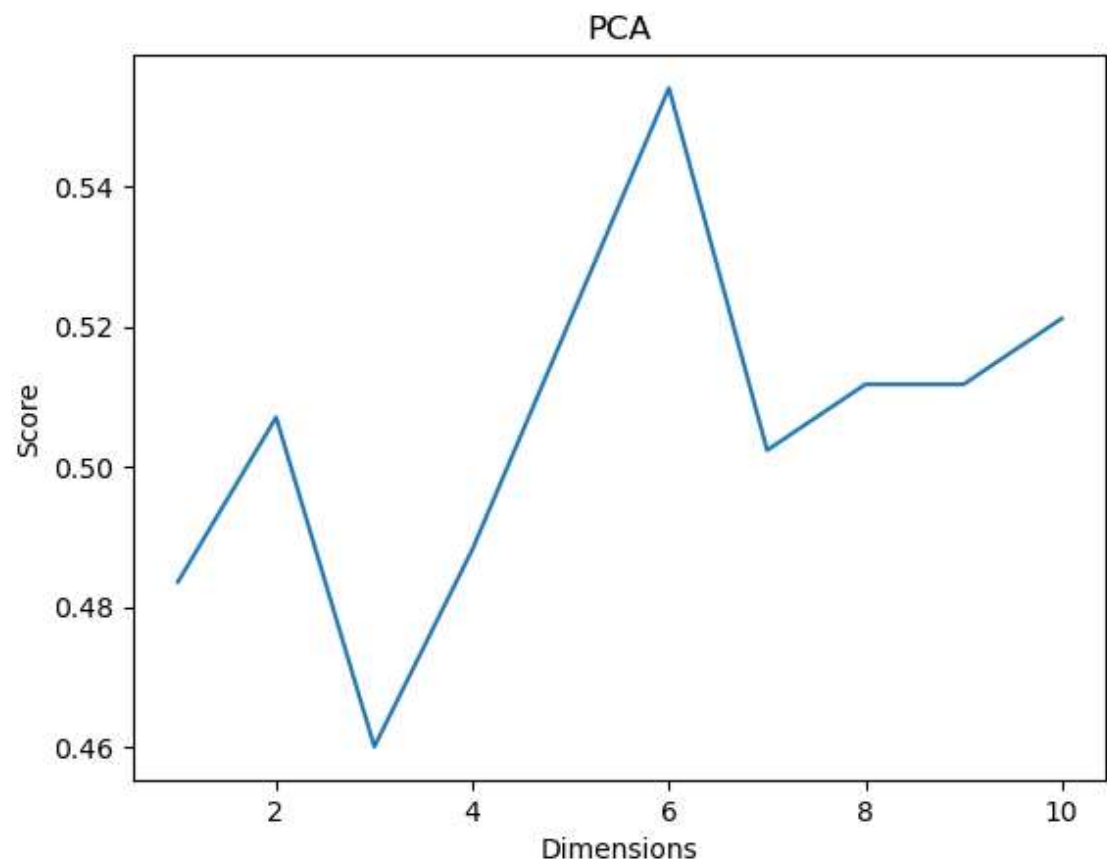
warnings.warn(

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

warnings.warn(

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

```
In [25]: ▶ plt.plot(list(range(1,11)), scores)
plt.title("PCA")
plt.xlabel("Dimensions")
plt.ylabel("Score")
plt.show()
```



```
In [26]: ▶ #LDA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
X = data.drop('target',axis = 1)
y = data['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
scores=[])

# X_Lda = Lda.fit(X, y).transform(X)
lda = LinearDiscriminantAnalysis(n_components=1)
X_train_reduced = lda.fit(X_train,y_train).transform(X_train)
lr = LogisticRegression()
lr.fit(X_train_reduced, y_train)
X_test_reduced = lda.transform(X_test)
scores.append(lr.score(X_test_reduced, y_test))
scores
```

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

warnings.warn(

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

warnings.warn(

C:\Users\Aaryan Agarwal\anaconda3\lib\site-packages\sklearn\utils\validation.py:1688: FutureWarning: Feature names only support names that are all strings. Got feature names with dtypes: ['int', 'str']. An error will be raised in 1.2.

warnings.warn(

Out[26]: [0.568075117370892]

Task 3

```
In [27]: ▶ df=pd.read_csv("./BCI-SSVEP_Database_Aceves/A002SB3_1.csv")
```

```
In [28]: df.head()
```

Out[28]:

	title:C	recorded:01.07.2012.54.57	sampling:128	subject:C	labels:COUNTER INTERPOLATED AF3 F7 F3 FC5 T7 P7 O1 O2 P8 T8 FC6 F4 F8 AF4	chan:37	unit
0	48	0	4226.666563	4219.487076	4215.897333	4215.897333	421:
1	49	0	4227.692204	4214.871692	4215.384512	4217.435794	422:
2	50	0	4230.769127	4224.102461	4216.410153	4223.076820	422:
3	51	0	4231.281948	4228.717845	4221.538358	4225.128102	421:
4	52	0	4226.666563	4219.999897	4215.897333	4218.461435	421:

In [29]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9600 entries, 0 to 9599
Data columns (total 16 columns):
 #   Column
Non-Null Count  Dtype
---  -
0   title:C
9600 non-null   int64
1   recorded:01.07.20 12.54.57
9600 non-null   int64
2   sampling:128
9600 non-null   float64
3   subject:C
9600 non-null   float64
4   labels:COUNTER INTERPOLATED AF3 F7 F3 FC5 T7 P7 O1 O2 P8 T8 FC6 F4
F8 AF4
9600 non-null   float64
5   chan:37
9600 non-null   float64
6   units:emotiv
9600 non-null   float64
7   Unnamed: 7
9600 non-null   float64
8   Unnamed: 8
9600 non-null   float64
9   Unnamed: 9
9600 non-null   float64
10  Unnamed: 10
9600 non-null   float64
11  Unnamed: 11
9600 non-null   float64
12  Unnamed: 12
9600 non-null   float64
13  Unnamed: 13
9600 non-null   float64
14  Unnamed: 14
9600 non-null   float64
15  Unnamed: 15
9600 non-null   float64
dtypes: float64(14), int64(2)
memory usage: 1.2 MB

```

In [30]: `df.corr()`

Out[30]:


	title:C	recorded:01.07.20 12.54.57	sampling:128	subject:C	labels:COUNTER INTERPOLATED AF3 F7 F3 FC5 T7 P7 O1 O2 P8 T8 FC6 F4 F8 AF4	
title:C	1.000000	NaN	0.033072	-0.013029	0.019828	-
recorded:01.07.20 12.54.57	NaN	NaN	NaN	NaN	NaN	
sampling:128	0.033072	NaN	1.000000	0.097258	-0.000708	-
subject:C	-0.013029	NaN	0.097258	1.000000	0.183393	
labels:COUNTER INTERPOLATED AF3 F7 F3 FC5 T7 P7 O1 O2 P8 T8 FC6 F4 F8 AF4	0.019828	NaN	-0.000708	0.183393	1.000000	
chan:37	-0.046208	NaN	-0.339962	0.110888	0.320856	
units:emotiv	0.023007	NaN	-0.010816	0.045856	0.161736	
Unnamed: 7	-0.030816	NaN	-0.076676	0.072644	0.275128	
Unnamed: 8	-0.033949	NaN	0.047691	0.032575	0.233974	
Unnamed: 9	-0.033426	NaN	0.013321	0.090028	0.004276	-
Unnamed: 10	-0.038339	NaN	0.044370	0.093762	0.254246	
Unnamed: 11	-0.030735	NaN	-0.090107	0.112440	0.058015	
Unnamed: 12	0.029623	NaN	-0.095490	-0.060605	0.428704	
Unnamed: 13	0.028798	NaN	0.074780	0.051832	0.480984	
Unnamed: 14	0.046023	NaN	0.002888	-0.004076	0.378384	
Unnamed: 15	-0.012715	NaN	0.206420	0.141856	0.183190	

In [31]: `import matplotlib.pyplot as plt`

In [32]: `x=df.iloc[81:210,0]
y=[df.iloc[81:210,3],df.iloc[81:210,4],df.iloc[81:210,5],df.iloc[81:210,6]`

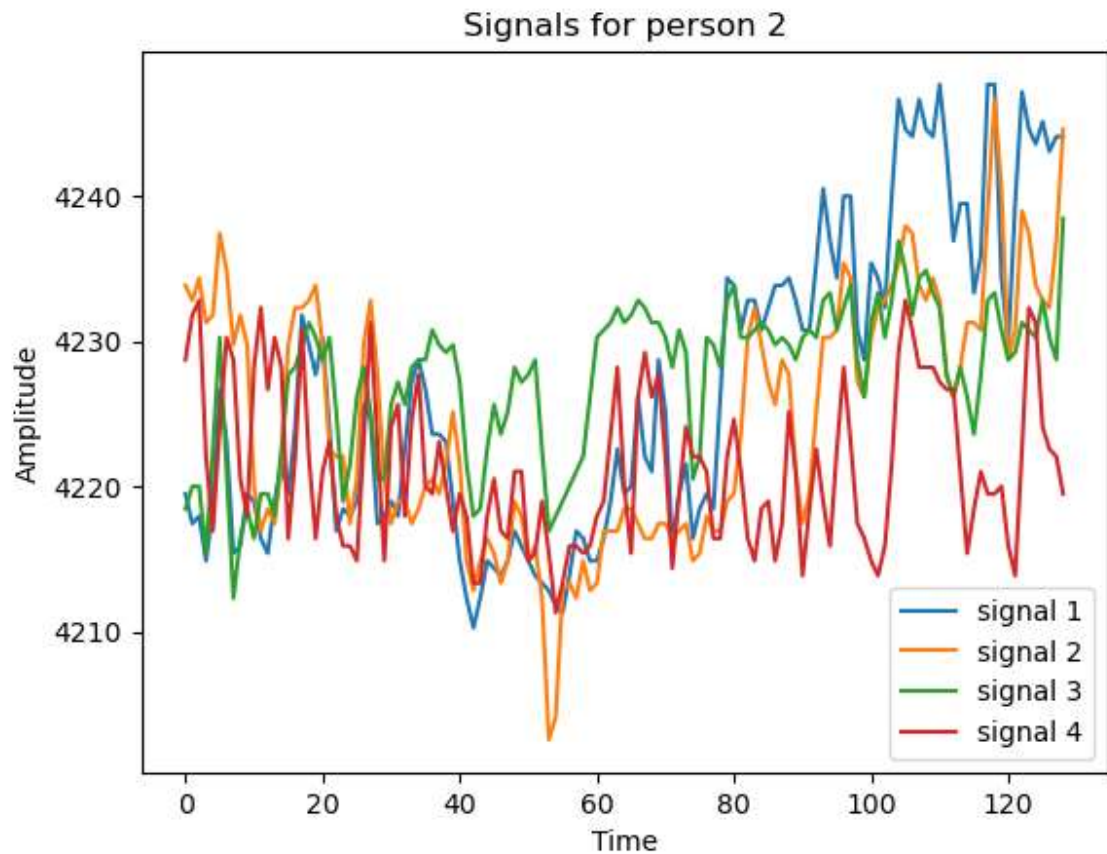
```
In [33]: ▶ plt.plot(x,y[0],label="signal 1")
plt.plot(x,y[1],label="signal 2")
plt.plot(x,y[2],label="signal 3")
plt.plot(x,y[3],label="signal 4")

plt.title("Signals for person 1")
plt.xlabel("Time")
plt.ylabel("Amplitude")
plt.legend()
plt.show()
# the graph shows the brain signal for person 1
```



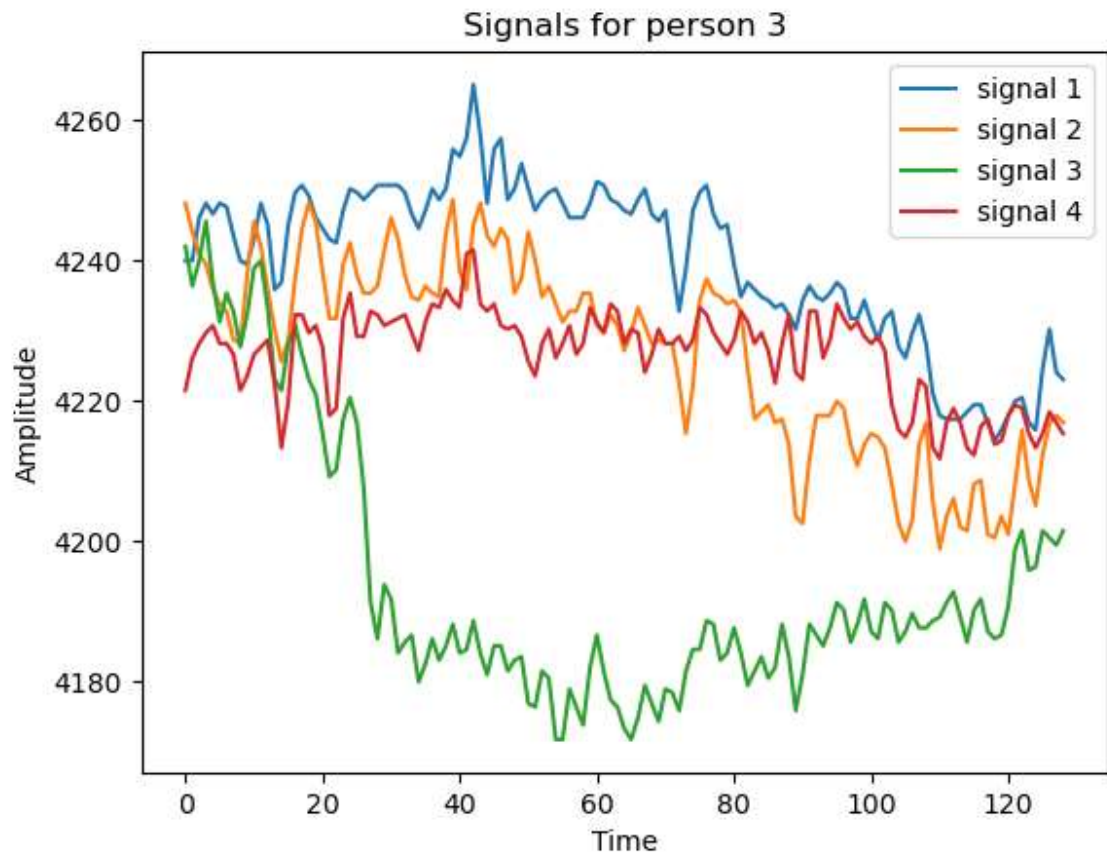
```
In [34]: x=df.iloc[210:339,0]
y=[df.iloc[210:339,3],df.iloc[210:339,4],df.iloc[210:339,5],df.iloc[210:339,6]]
plt.plot(x,y[0],label="signal 1")
plt.plot(x,y[1],label="signal 2")
plt.plot(x,y[2],label="signal 3")
plt.plot(x,y[3],label="signal 4")

plt.title("Signals for person 2")
plt.xlabel("Time")
plt.ylabel("Amplitude")
plt.legend()
plt.show()
# the graph shows the brain signals from person 2
```



```
In [35]: x=df.iloc[339:468,0]
y=[df.iloc[339:468,3],df.iloc[339:468,4],df.iloc[339:468,5],df.iloc[339:468,6]]
plt.plot(x,y[0],label="signal 1")
plt.plot(x,y[1],label="signal 2")
plt.plot(x,y[2],label="signal 3")
plt.plot(x,y[3],label="signal 4")

plt.title("Signals for person 3")
plt.xlabel("Time")
plt.ylabel("Amplitude")
plt.legend()
plt.show()
# the graph shows the brain signals for person 3
```



```
In [36]: df=pd.read_csv("../repeat_consumption_data/data/reddit_sample/test.csv",hea
```

In [37]: `df.head()`

Out[37]:

	0	1	2
0	1	575	4
1	1	3063	2
2	1	5576	3
3	1	11561	1
4	1	11835	2

In [38]: `df.corr()`

Out[38]:

	0	1	2
0	1.000000	0.002036	-0.011609
1	0.002036	1.000000	-0.015056
2	-0.011609	-0.015056	1.000000

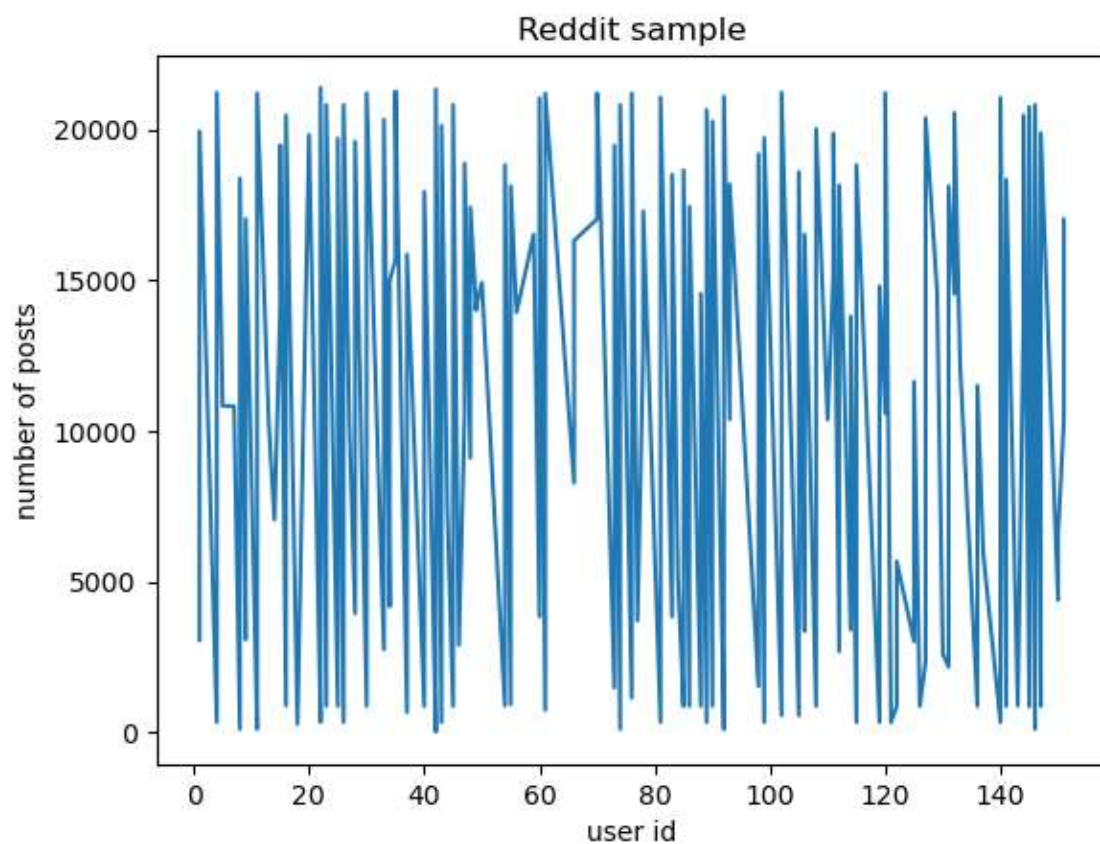
In [39]: `# for i in range(10):
signal = df.iloc[1:1000]
print(signal)`

	0	1	2
1	1	3063	2
2	1	5576	3
3	1	11561	1
4	1	11835	2
5	1	12396	1
..
995	150	5947	1
996	150	5958	1
997	150	6569	6
998	151	10361	1
999	151	17016	3

[999 rows x 3 columns]

In [40]: `x=signal[0]
y=signal[1]`

```
In [41]: ▶ plt.plot(x,y)
plt.title("Reddit sample")
plt.xlabel("user id")
plt.ylabel("number of posts")
plt.show()
```



```
In [ ]: ▶ # the number of posts made by each user id
```

```
In [ ]: ▶
```